

VM-Series for GCP



GCP Terraform Template Deployment Guide

Deploys an External Load Balancer multiple VM-Series NGFW Internal Load Balancers and Web Servers in different zone. This deployment model is commonly referred to as a Load Balancer Sandwich.

<https://www.paloaltonetworks.com>

Table of Contents

Version History.....	3
1. About Terraform Templates.....	4
2. Support Policy.....	5
3. Instances used	6
4. Prerequisites	6
4.1 Create GCP account.....	6
4.2 Install the Google Cloud SDK	6
4.3 Accept the EULA (If Required).....	6
4.4 Create a Project.....	6
4.5 Enable the API	8
4.6 Create a Bootstrap Bucket.....	9
4.7 Download the Terraform Template Files.....	12
4.8 Extract the Files	Error! Bookmark not defined.
4.9 Gather Information and Update the Template File	12
5. Launch the Template	13
6. Review what was created.....	15
7. Access the firewall.....	20
8. Access the Webservers via ELB	25
9. Cleanup	25
9.1 Delete the deployment	25
10. Conclusion	26
Appendix A	26
Troubleshooting tips.....	26

Version History

Version number	Comments
1.0	Initial Draft

1. About Terraform Templates

GCP Terraform Templates, are files that can deploy, configure, and launch GCP resources such as VPC networks & subnets, security groups, firewall rules, route tables, load balancers, and more. These templates are used for ease of deployment and are key to any cloud deployment model.

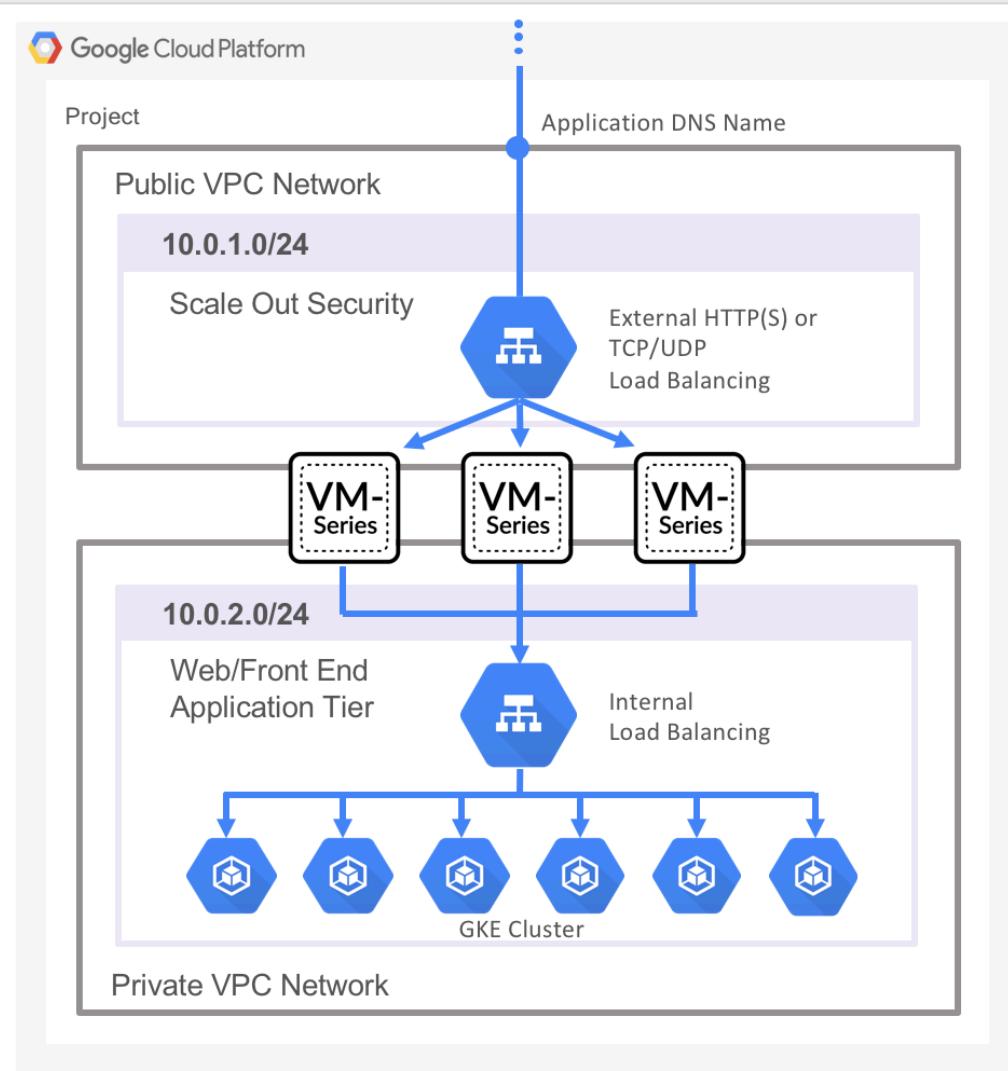
For more information on Templates refer to Google's documentation

<https://cloud.google.com/community/tutorials/managing-gcp-projects-with-terraform>

There are also many Terraform template s available here:

<https://github.com/GoogleCloudPlatform/terraform-google-examples>

This document will explain how to deploy a Terraform template that launches everything that is shown below in the diagram. This includes, multiple apache web server, multiple VM-Series firewall and the subnets, an HTTP ELB, and a TCP ILB. In addition, the Terraform template performs a native bootstrapping feature on the VM-Series firewall that allows for additional configuration of the VM-Series firewall (such as routes, security policies, management interface swap, etc.) Once the Terraform template has been deployed, the network topology will align with the following diagram:



2. Support Policy

This template is released under an as-is, best effort, support policy. These scripts should be seen as community supported and Palo Alto Networks will contribute our expertise as and when possible. We do not provide technical support or help in using or troubleshooting the components of the project through our normal support options such as Palo Alto Networks support teams, or ASC (Authorized Support Centers) partners and backline support options. The underlying product used (the VM-Series firewall) by the scripts or templates are still supported, but the support is only for the product functionality and not for help in deploying or using the template or script itself.

Unless explicitly tagged, all projects or work posted in our GitHub repository (at <https://github.com/PaloAltoNetworks/googlecloud>) or sites other than our official Downloads page on <https://support.paloaltonetworks.com> are provided under the best effort policy.

3. Instances used

When using this Terraform template the following machine types are used:

Instances	Machine Types
Apache Web Servers	f1-micro
VM Series Firewall	n1-standard-4

Note: There are costs associated with each machine type launched, please refer to the Google instance pricing page <https://cloud.google.com/compute/pricing>

4. Prerequisites

Here are the prerequisites required to successfully launch this template:

- Terraform installed

4.1 Create GCP account

If you do not have a GCP account already, go to <https://cloud.google.com/free/> and create an account.

4.2 Install the Google Cloud SDK

Template installations in GCP are performed from the CLI. Install the SDK/CLI by selecting the relevant platform from the following link and following the installation instructions:

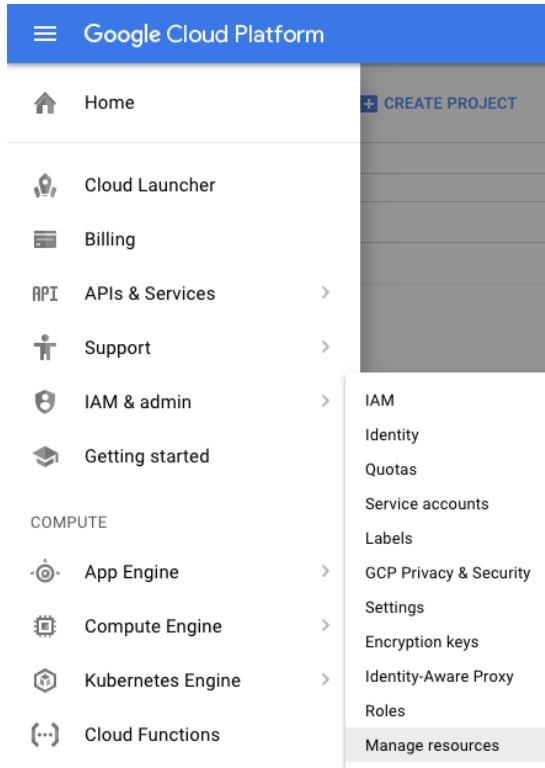
<https://cloud.google.com/sdk/>

4.3 Accept the EULA (If Required)

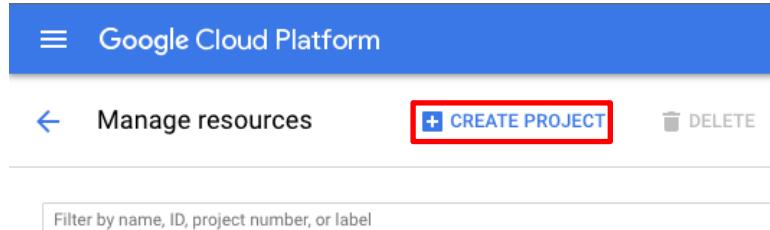
4.4 Create a Project

All GCP resources are deployed to a GCP Project. A GCP Project is an organizational boundary that separates users, resources, billing information, etc. A GCP Project is similar to an AWS VPC or an Azure Resource Group. By default, GCP will create a Project upon creation of an account. If that is not the case or to manually create a dedicated project, use the drop-down on the left and select **IAM & admin > Manage Resources**:

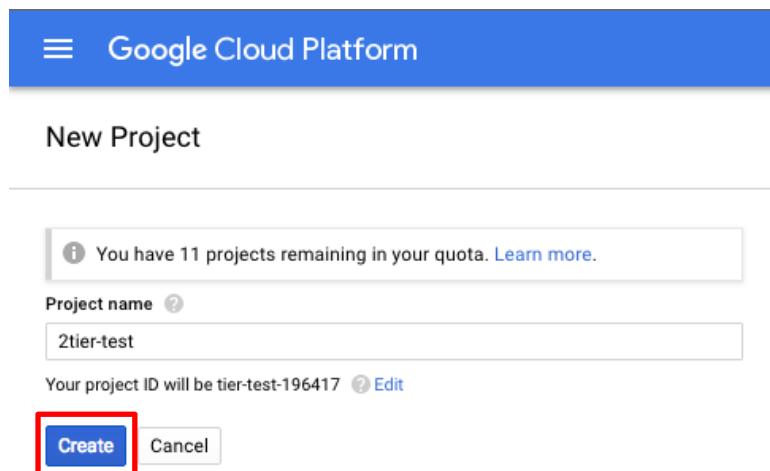
Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich



Click **Create Project**:



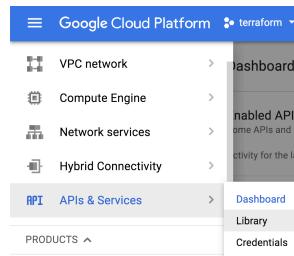
Specify a name for the project and click **Create**:



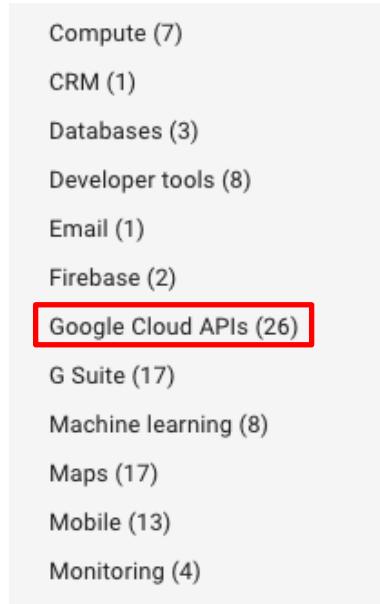
Note: GCP Project creation will take a few minutes.

4.5 Enable the API

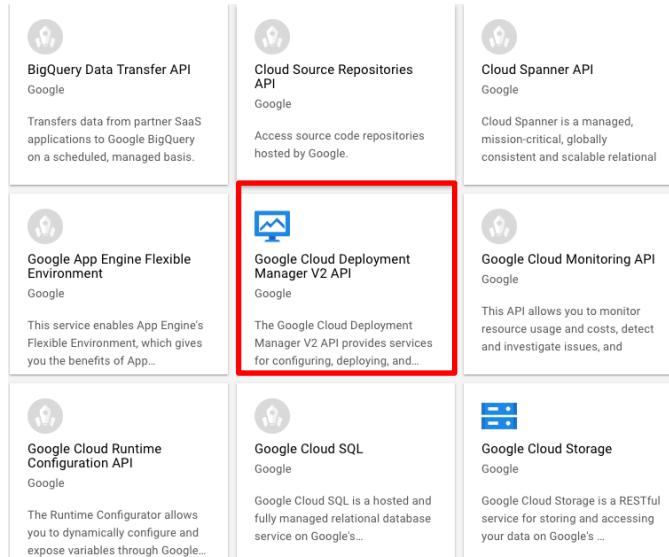
Deploying a template requires the API be enable on the project. Navigate to **APIs & Services > Library**:



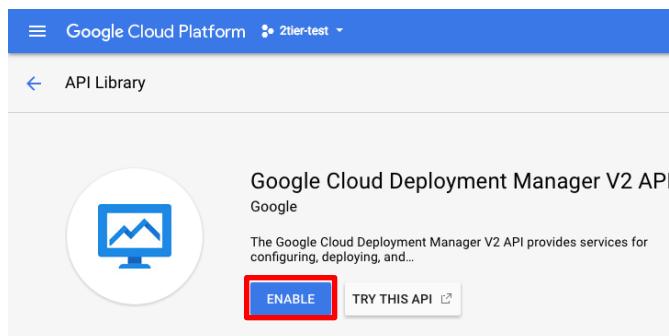
Select Google Cloud APIs on the left-hand-side:



Select **Google Cloud Deployment Manager V2 API**:



Select **Enable**:

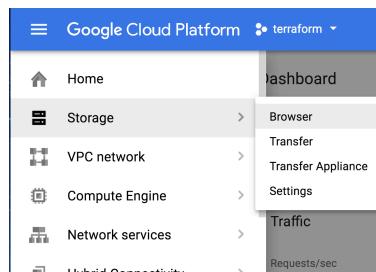


Note: Enabling the API for the project will take a few minutes to complete.

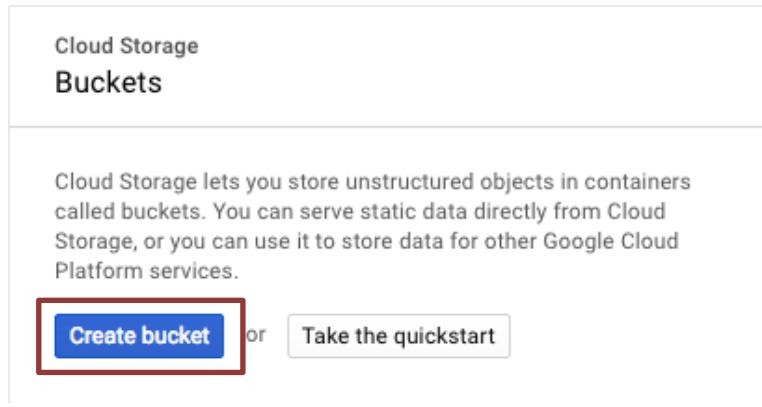
4.6 Create a Bootstrap Bucket

Bootstrapping is a feature of the VM-Series firewall that allows you to load a pre-defined configuration into the firewall during boot-up. This ensures that the firewall is configured and ready at initial boot-up, thereby removing the need for manual configuration. The bootstrapping feature also enables automating deployment of the VM-Series firewall.

In order to create a Bootstrap bucket, navigate to **Storage > Browser**:



Click **Create Bucket**:



Specify a globally-unique bucket name and regional settings and click **Create**:

The screenshot shows the 'Create a bucket' dialog. It has a back arrow and the title 'Create a bucket'. Under 'Name', it says 'Must be unique across Cloud Storage. If you're [serving website content](#), enter the website domain as the name.' A text input field contains 'bootstrap-elb'. Under 'Default storage class', it says 'Compare storage classes' and shows 'Multi-Regional' selected (radio button is checked). Other options are 'Regional', 'Nearline', and 'Coldline'. Under 'Location', a dropdown menu shows 'United States'. At the bottom, there are four cost breakdowns: 'Storage cost \$0.026 per GB-month', 'Retrieval cost Free', 'Class A operations \$0.005 per 1,000 ops', and 'Class B operations \$0.0004 per 1,000 ops'. At the very bottom are 'Create' and 'Cancel' buttons.

You will need to enter a globally unique bucket name. GCP will warn you if the name is not unique. Once the bucket is created, click on the newly created bucket and add four folders called **config**, **license**, **software** and **content** by clicking on **Create Folder**:

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

The screenshot shows a GCP Storage Bucket interface. At the top, there are buttons for 'UPLOAD FILES' (blue), 'CREATE FOLDER' (red box), 'REFRESH' (blue), 'SHARE PUBLICLY' (grey), and 'DELETE' (grey). Below the buttons is a search bar with placeholder text 'Filter by prefix...'. Underneath the search bar, it says 'Buckets / 2tier-bootstrap'. A table lists five items:

Name	Size	Type	Storage class	Last modified
config/	—	Folder	—	—
content/	—	Folder	—	—
license/	—	Folder	—	—
software/	—	Folder	—	—

Download the following files using the links provided and save the files in a known location:

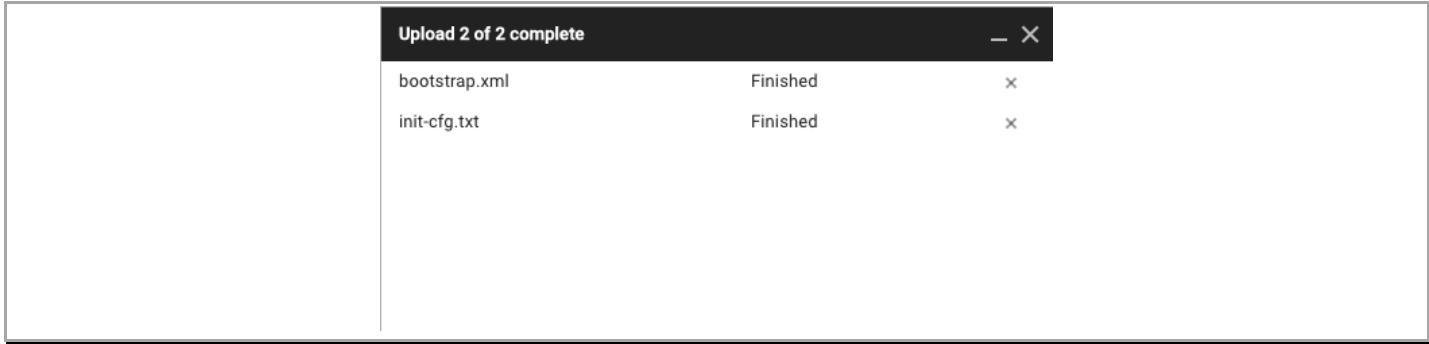
<https://github.com/PaloAltoNetworks/GCP-Terraform-Samples/blob/master/LB%20Sandwich/bootstrap.zip>

<https://github.com/PaloAltoNetworks/GCP-Terraform-Samples/blob/master/LB%20Sandwich/startup.sh>

Now click on the **config** folder in the console and click **UPLOAD FILES**:

The screenshot shows a GCP Storage Bucket interface. The 'config/' folder is selected (highlighted with a red box). At the bottom, there is a horizontal bar with three buttons: 'Browser' (grey), 'UPLOAD FILES' (red box), and 'UPLOAD FOLDER' (blue).

Select the two files (bootstrap.xml and init-cft.txt) downloaded previously and click **Open**:



NOTE: All four folders must be created for the bootstrapping process to occur. However, all folders DO NOT need to contain files.

NOTE: Please create the folders using the GUI or GCP CLI console. Creating folders locally on your machine and uploading them may not work as expected.

4.7 Download the Terraform Template Files

Download and save all of the template files to a known location by selecting **Clone or download**:

<https://github.com/PaloAltoNetworks/GCP-Terraform-Samples/blob/master/LB%20Sandwich/Main.tf>

<https://github.com/PaloAltoNetworks/GCP-Terraform-Samples/blob/master/LB%20Sandwich/Variables.tf>

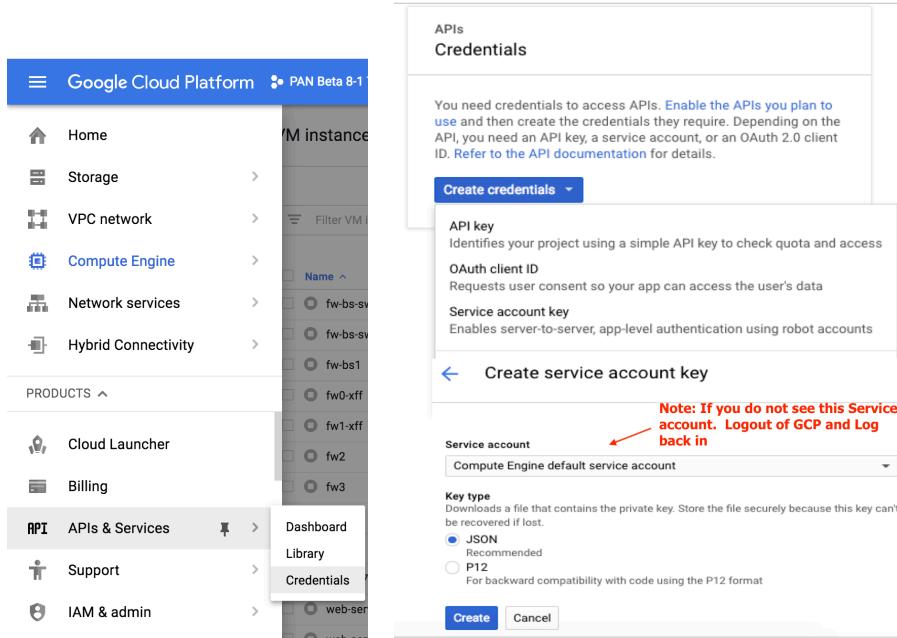
<https://github.com/PaloAltoNetworks/GCP-Terraform-Samples/blob/master/LB%20Sandwich/Output.tf>

4.8 Gather Information and Update the Template File

Deploying the Terraform template in GCP requires modification of the template Main and Variable files to include deployment-specific information. The minimum required information is:

```
credentials = "${file("Your_Project_Credentials.json")}"  
  
metadata_startup_script= "${file("location/startup.sh")}"  
  
project    = "${Your_Project_ID}"  
  
region     = "${Your_Project_Region}"  
  
zone       = "${Your_Project_Zone}"  
  
2nd zone    = "${Your_Zone2}"  
  
sshKey     = "${Your_Public_SSHKey}"
```

To create the credentials to access the APIs in JSON format. In GCP console go to (APIs & Services > Credentials > Create Credentials > Service Account Key), and download the file (client_secrets.json). Put the .json credential file in your Terraform template folder.



Once the information has been gathered, update the Main and Variable files with the information. Save the Files.

5. Launch the Template

Navigate to a command shell navigate to the directory containing the downloaded template files:

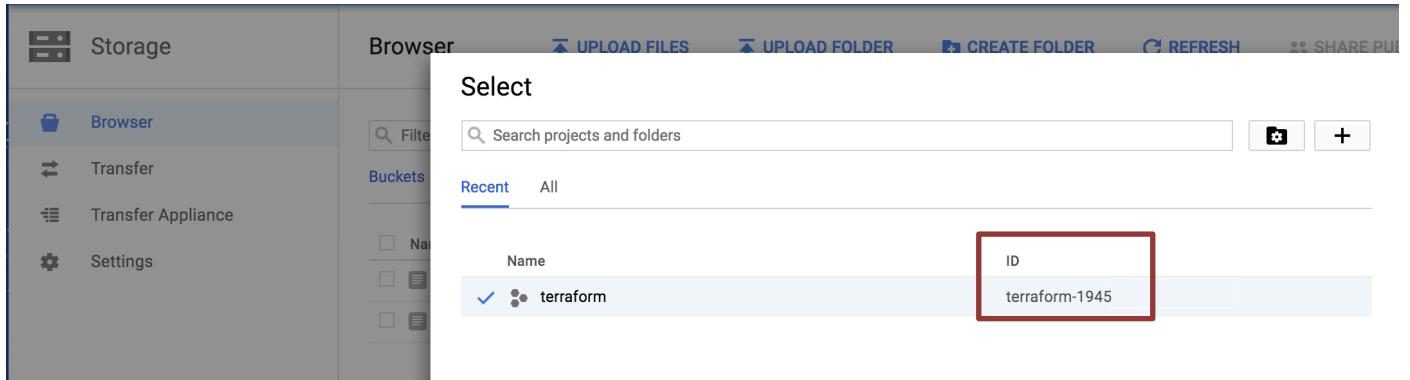
Authenticate to the GCP environment from the command line with the command:

```
$ gcloud auth login
```

- Copy/paste the link into a browser and select the account to authenticate if a browser does not automatically launch:
- Review the requested permissions and click **Allow**:
- Copy the one-time verification code:
- Paste it into the window to complete the authentication request (ignore the warning):

Get the Project ID:

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich



Set the target project for template deployment via command line:

```
$ gcloud config set project my_Project_id
```

Run Terraform Commands:

Initiate template deployment using command “**terraform init**”.

```
Terraform has been successfully initialized!
You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

Once the terraform init has completed run the command “**terraform plan**”.

```
Plan: 12 to add, 0 to change, 0 to destroy.

-----
Note: You didn't specify an "-out" parameter to save this plan, so Terraform
can't guarantee that exactly these actions will be performed if
"terraform apply" is subsequently run.
```

You will see if there are any errors and what terraform will be deploying. Now run the “**terraform apply**” command and say **yes** when prompt.

```
Plan: 12 to add, 0 to change, 0 to destroy.
Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

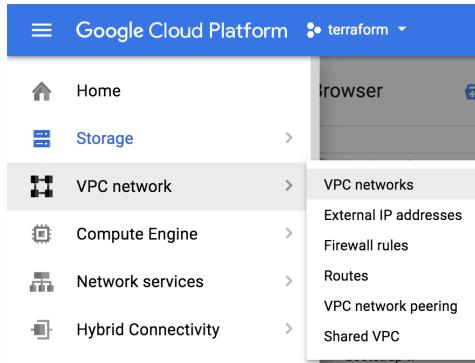
Enter a value: yes
```

If all goes well, Terraform will report success (“Apply Complete!” and no errors):

```
Apply complete! Resources: 12 added, 0 changed, 0 destroyed.
```

6. Review what was created

Let's review what the template has launched. The newly created networks can be viewed via **VPC Networks > VPC Network**:



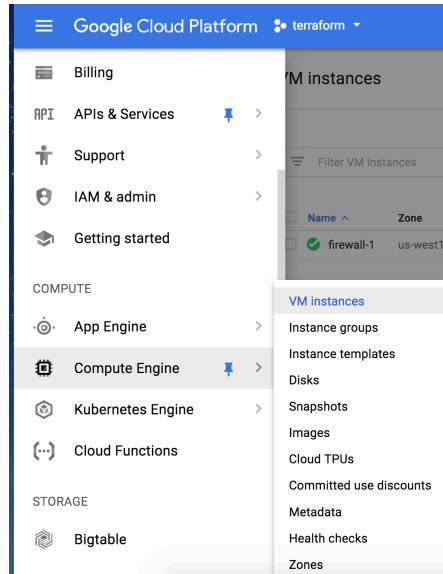
The template creates four networks: management-network, trust-network, and untrust-network.

VPC networks							
Name	Region	Subnets	Mode	IP addresses ranges	Gateways	Firewall Rules	Global dynamic routing
management	us-west1	management-sub	Custom	10.0.0.0/24	10.0.0.1	1	Off
trust	us-west1	trust-sub	Custom	10.0.2.0/24	10.0.2.1	1	Off
untrust	us-west1	untrust-sub	Custom	10.0.1.0/24	10.0.1.1	1	Off

Note: A default network is automatically created when a GCP Project is instantiated. This default network can be ignored or deleted.

Deployed hosts can be viewed by navigating to **Compute Engine > VM Instances**:

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich



High-level information regarding the deployed instances are available with the default view:

The screenshot shows the 'VM instances' list with the following data:

Name	Zone	Recommendation	Internal IP	External IP	Connect
firewall-1	us-west1-a		10.0.1.2	35.230.119.56	SSH
firewall-2	us-west1-a		10.0.1.3	35.230.69.250	SSH
firewall-3	us-west1-a		10.0.1.4	35.185.245.54	SSH
webserver-1	us-west1-a		10.0.2.3	None	SSH
webserver-2	us-west1-b		10.0.2.7	None	SSH
webserver-3	us-west1-a		10.0.2.6	None	SSH
webserver-4	us-west1-b		10.0.2.8	None	SSH

Also note the order in which the networks are attached to the firewalls. Click on firewall-1 and scroll down to see the network order.

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

VM instance details EDIT RESET CLONE STOP DELETE

firewall-1

CPU utilization

CPU

% CPU

Mar 6, 10:00 AM Mar 6, 10:05 AM Mar 6, 10:10 AM Mar 6, 10:15 AM Mar 6, 10:20 AM Mar 6, 10:25 AM

CPU: 6.472

Remote access

SSH Connect to serial console

Enable connecting to serial ports

Logs

Stackdriver Logging

Serial port 1 (console)

More

Machine type

n1-standard-4 (4 vCPUs, 15 GB memory)

In use by

fw-ig

CPU platform

Intel Broadwell

Zone

us-west1-a

Labels

None

Creation time

Mar 6, 2018, 9:44:20 AM

Network interfaces					
Network	Subnetwork	Primary internal IP	Alias IP ranges	External IP	IP forwarding
untrust	untrust-sub	10.0.1.2	—	35.230.119.56 (ephemeral)	On
management	management-sub	10.0.0.2	—	35.185.216.43 (ephemeral)	
trust	trust-sub	10.0.2.2	—	None	

Public DNS PTR Record

None

Network tags

None

NOTE: The untrust network is first. The GCP Load Balancers only communicate with the lowest numbered interface on a VM. During the bootstrap phase of deployment the init-cfg.txt told the VM-Series firewall to perform a management interface swap. Therefore, we must have the GCP networks in this order.

In order for the webservers to download apache we have added a default route through the first VM-Series Firewall instance. Navigate to routes via the VPN Network dropdown. See below.

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

The screenshot shows the navigation menu on the left with 'VPC network' selected. A sub-menu for 'Routes' is open, showing options like 'Routes', 'VPC network peering', and 'Shared VPC'. Below this, the 'Routes' section is displayed with a table of routes. The table has columns: Name, Destination IP ranges, Priority, Instance tags, Next hop, and Network. The routes listed are:

Name	Destination IP ranges	Priority	Instance tags	Next hop	Network
default-route-354346117ff5059a	10.0.0.0/24	1000	None	Virtual network	management
default-route-4ec4d794bfa971a0	10.0.2.0/24	1000	None	Virtual network	trust
default-route-6726ace17905c873	0.0.0.0/0	1000	None	Default internet gateway	untrust
default-route-7832e335985081ec	10.0.1.0/24	1000	None	Virtual network	untrust
default-route-85c1ffd65b473093	0.0.0.0/0	1000	None	Default internet gateway	trust
default-route-dde47d271469593f	0.0.0.0/0	1000	None	Default internet gateway	management
trust-route	0.0.0.0/0	100	None	firewall-1 (Zone us-west1-a)	trust

Lastly check your newly deployed Load Balancers by navigating to Network Services then select Load Balancing.

The screenshot shows the Google Cloud Platform navigation menu with 'Network services' selected. A sub-menu for 'Load balancing' is open, showing options like 'Cloud DNS', 'Cloud CDN', and 'Networking solutions'. The 'Load balancing' option is highlighted.

Load balancing

[CREATE LOAD BALANCER](#)

[REFRESH](#)

[DELETE](#)

Load balancers

Backends

Frontends



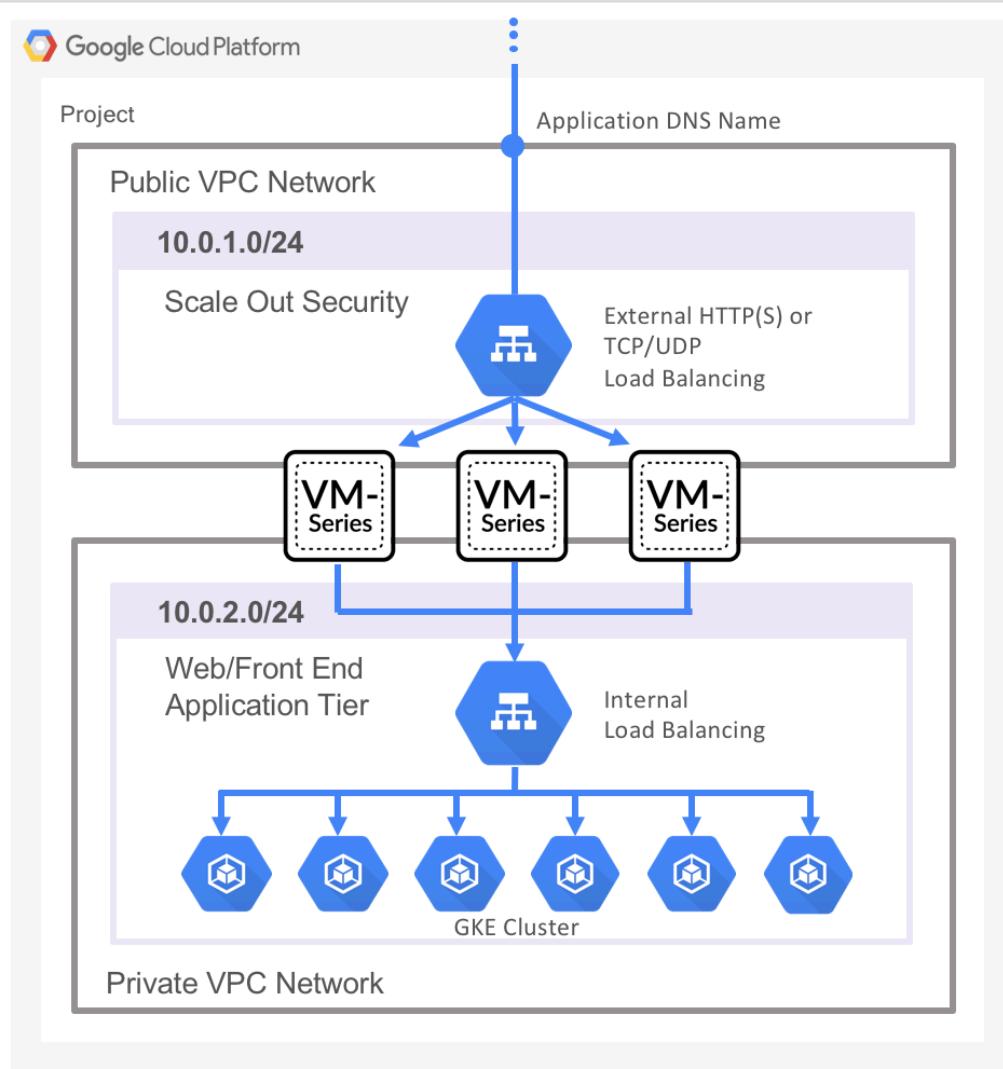
Filter by name or protocol

<input type="checkbox"/> Name	Protocol ^	Backends	
<input type="checkbox"/> http-elb	HTTP	1 backend service is unhealthy (group)	
<input type="checkbox"/> my-int-lb	TCP (Internal)	1 regional backend service (2 instance groups)	

To edit load balancing resources like forwarding rules and target proxies, go to the [advanced menu](#).

NOTE: You should see two load balancers configured with the HTTP-ELB showing that the backend service is unhealthy. The Internal TCP LB shows all backend services as healthy. We need create NAT rules on the VM-Series Firewall to achieve a healthy state for the HTTP-ELB backend services in the next session.

All of this matches the topology shown previously:



7. Access the firewall

NOTE: Bootstrapping a VM-Series firewall takes approximately 9 minutes. Be patient ☺ Once the template has been deployed successfully, it may be a while before the VM-Series firewall is up and you are able to log into the VM-Series firewall by browsing to the Management public IP Address. Recall we swapped the Management interface so you will need to click on the VM Series to get the Public IP address.

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

VM instance details EDIT RESET CLONE STOP DELETE

firewall-1

CPU utilization ▾

CPU

% CPU

The chart displays CPU usage in percent over a period from March 6, 10:30 AM to 10:50 AM. The Y-axis ranges from 0 to 8%. The X-axis shows five time points: Mar 6, 10:30 AM, Mar 6, 10:35 AM, Mar 6, 10:40 AM, Mar 6, 10:45 AM, and Mar 6, 10:50 AM. The usage fluctuates between approximately 5% and 7%.

Mar 6, 10:30 AM Mar 6, 10:35 AM Mar 6, 10:40 AM Mar 6, 10:45 AM Mar 6, 10:50 AM

CPU: 5.449

Remote access

SSH ▾ Connect to serial console ▾

Enable connecting to serial ports ⓘ

Logs

Stackdriver Logging
Serial port 1 (console)
More

Machine type

n1-standard-4 (4 vCPUs, 15 GB memory)

In use by

fw-ig

CPU platform

Intel Broadwell

Zone

us-west1-a

Labels

None

Creation time

Mar 6, 2018, 9:44:20 AM

Network interfaces

Network	Subnetwork	Primary internal IP	Alias IP ranges	External IP	IP forwarding
untrust	untrust-sub	10.0.1.2	—	35.230.119.56 (ephemeral)	On
management	management-sub	10.0.0.2	—	35.185.216.43 (ephemeral)	
trust	trust-sub	10.0.2.2	—	None	

Public DNS PTR Record

None

Network tags

None

You should now be able to browse to the VM-Series firewall and login using the **username: paloalto** and password: **Pal0Alt0@123**

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

The screenshot shows the Palo Alto Networks Management Console dashboard. It includes sections for General Information, Logged In Admins, System Logs, and System Resources. The General Information section displays various device details like Device Name, MGT IP Address, and Software Version. The Logged In Admins section shows an admin session from 12.206.19.5. The System Logs section contains a log entry for failed password attempts from 192.169.155.230. The System Resources section shows CPU usage and session counts.

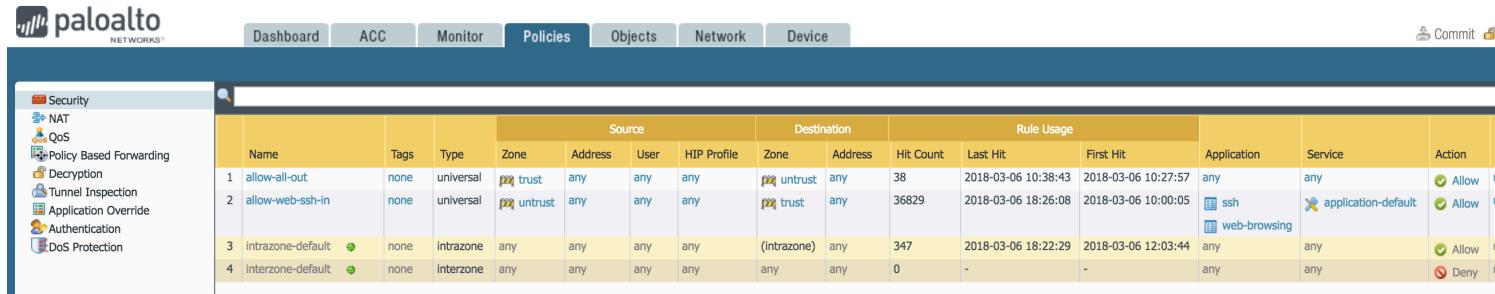
Here are the interfaces to zone mappings.

The screenshot shows the Palo Alto Networks Management Console Network tab. The left sidebar lists various network-related configurations. The main pane displays a table of interface-to-zone mappings. The table has columns for Interface, Interface Type, Management Profile, Link State, IP Address, Virtual Router, Tag, VLAN / Virtual-Wire, Security Zone, Features, and Comment. Several rows in the table show different interfaces (e.g., ethernet1/1, ethernet1/2) mapped to security zones like 'untrust' and 'trust'. A red box highlights the 'Security Zone' column for these rows.

Interface	Interface Type	Management Profile	Link State	IP Address	Virtual Router	Tag	VLAN / Virtual-Wire	Security Zone	Features	Comment
ethernet1/1	Layer3	mgmt-untrust	Dynamic-DHCP Client	default		Untagged	none	untrust		
ethernet1/2	Layer3	mgmt-trust	Dynamic-DHCP Client	default		Untagged	none	trust		
ethernet1/3			none	none		Untagged	none	none		
ethernet1/4			none	none		Untagged	none	none		
ethernet1/5			none	none		Untagged	none	none		
ethernet1/6			none	none		Untagged	none	none		
ethernet1/7			none	none		Untagged	none	none		

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

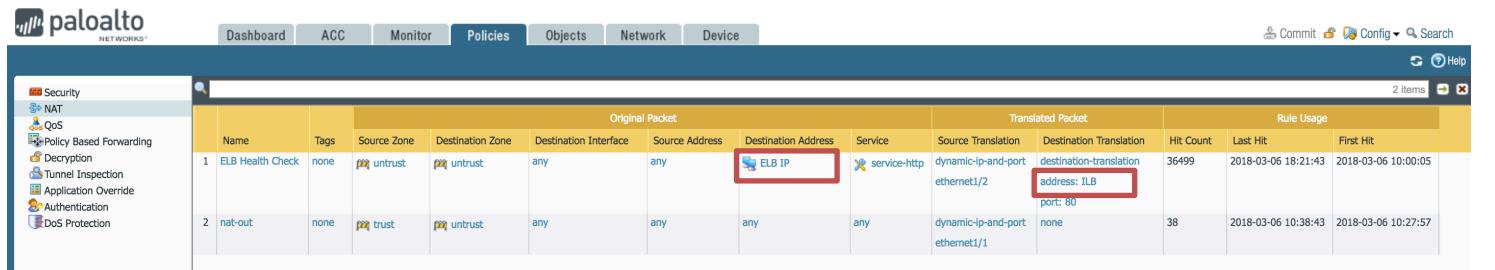
In the policies tab you can review the security policies:



The screenshot shows the Palo Alto Networks Firewall UI with the 'Policies' tab selected. On the left sidebar, under 'Security', the 'NAT' option is highlighted. The main pane displays a table of security rules:

Name	Tags	Type	Source			Destination			Rule Usage				
			Zone	Address	User	HIP Profile	Zone	Address	Hit Count	Last Hit	First Hit	Application	Action
1 allow-all-out	none	universal	trust	any	any	any	untrust	any	38	2018-03-06 10:38:43	2018-03-06 10:27:57	any	any
2 allow-web-ssh-in	none	universal	untrust	any	any	any	trust	any	36829	2018-03-06 18:26:08	2018-03-06 10:00:05	ssh	application-default
3 intrazone-default	none	intrazone	any	any	any	any	(intrazone)	any	347	2018-03-06 18:22:29	2018-03-06 12:03:44	any	any
4 interzone-default	none	interzone	any	any	any	any	any	any	0	-	-	any	any

Under the polices tab select NAT on the left side. We need to have a health check from the ELB through to the backend web server of the ILB. In order for the ELB health checks to flow through the VM-Series firewall then to and through the ILB to the web servers we need to add the Terraform outputs to the highlighted objects in the ELB Health Check NAT statement in the next step. There is also a rule to NAT all traffic out from web servers to the outside world.



The screenshot shows the Palo Alto Networks Firewall UI with the 'NAT' tab selected. On the left sidebar, under 'Security', the 'NAT' option is highlighted. The main pane displays a table of NAT rules:

Name	Tags	Source Zone	Destination Zone	Destination Interface	Original Packet			Translated Packet			Rule Usage			
					Source Address	Destination Address	Service	Source Translation	Destination Translation	Hit Count	Last Hit	First Hit		
1 ELB Health Check	none	untrust	untrust	any	any	ELB IP	service-https	dynamic-ip-and-port	destination-translation	36499	2018-03-06 18:21:43	2018-03-06 10:00:05		
2 nat-out	none	trust	untrust	any	any	any	any	dynamic-ip-and-port	address: ILB port: 80	38	2018-03-06 10:38:43	2018-03-06 10:27:57		

Make note of the outputs from your Terraform apply. Your outputs may be different.

```
firewall-untrust-ips-for-nat-healthcheck = [
  10.0.1.2,
  10.0.1.3,
  10.0.1.4
]
internal-lb-ip-for-nat-healthcheck = 10.0.2.9
```

Click on the Objects Tab and up the ELB IP and ILB on each VM-Series Firewall.

Palo Alto Networks GCP Terraform Template Deployment Guide LB Sandwich

The screenshot shows the Palo Alto Networks Firewall interface. The top navigation bar includes tabs for Dashboard, ACC, Monitor, Policies, Objects (which is selected), Network, and Device. On the left, a sidebar menu lists various objects: Addresses, Address Groups, Regions, Applications, Application Groups, Application Filters, Services, Service Groups, and Tags. The main content area displays a table titled 'Addresses' with columns for Name, Location, Type, Address, and Tags. Two entries are listed: 'ELB IP' and 'ILB', both of which are IP Netmask types with addresses 10.0.1.2 and 10.0.2.9 respectively.

Repeat this step for all of the Firewalls and **COMMIT** these changes on the VM-Series Firewalls
Then go back to your GCP Load Balancers and check the Health Status.

The screenshot shows the GCP Load Balancer configuration for a load balancer named 'http-eb'. The interface is divided into several sections:

- Frontend:** Shows a single rule for HTTP on port 80, mapping to an internal IP of 35.227.197.98:80.
- Host and path rules:** Shows a default rule for 'All unmatched (default)' hosts and paths, pointing to the 'fw-backend' backend.
- Backend:** Shows the 'fw-backend' service with the following details:
 - Endpoint protocol: HTTP
 - Named port: http
 - Timeout: 30 seconds
 - Health check: elb-health-check
 - Session affinity: None
 - Cloud CDN: disabled
- Backend services:** A detailed view of the 'fw-backend' service, including its configurations and health status.
- 1. fw-backend:** Summary of the service with the following parameters:
 - Endpoint protocol: HTTP
 - Named port: http
 - Timeout: 30 seconds
 - Health check: elb-health-check
 - Session affinity: None
 - Cloud CDN: disabled
- Advanced configurations:** A collapsed section under the service summary.
- Instance group:** A table showing the健康 (Healthy) status of the instance group 'fw-ig' in zone 'us-west1-a'. The table has columns for Instance group, Zone, Healthy, Autoscaling, Balancing mode, and Capacity. The 'Healthy' column for 'fw-ig' shows '3 / 3' and is highlighted with a red box.

8. Access the Webservers via ELB

Open a browser and browse to the IP address of the ELB. The IP of the ELB can be found under load balancers then expand the ELB. Each Webserver will provide their name in the browser upon reloading the browser. This lets you know the Load Balancer is working.

The screenshot shows the Palo Alto Networks Cloud UI interface. On the left, there's a sidebar with icons for Load balancing, Cloud DNS, and Cloud CDN. The main area has a header with tabs: Load balancers (which is selected), Backends, and Frontends. Below the header, it says "Load balancer" and "Protocol". There's a list item "http-elb" with a green checkmark. Under "http-elb", there are three tabs: Details (selected), Monitoring, and Caching. In the "Frontend" section, there's a table with one row for "HTTP" where the "IP:Port" field contains "35.227.197.98:80". This entire row is highlighted with a red box. Below the frontend, there's a section for "Host and path rules" with tabs for Hosts, Paths, and Backend. The "Hosts" tab is selected, showing "All unmatched (default)". The "Backend" tab shows "fw-backend". At the bottom, there's a section for "Backend".

You have now successfully deployed a Terraform template with a VM-Series firewall in GCP.

9. Cleanup

9.1 Delete the deployment

Once done, cleanup as follows:

- If you licensed the VM-Series firewall perform the De-License function.
 - https://www.paloaltonetworks.com/documentation/71/virtualization/virtualization/license-the-vm-series-firewall/deactivate-vm#_87329
- From the CLI, issue the command “**terraform destroy**”
 - This will delete all the resources created via the Terraform template.

10. Conclusion

You have successfully deployed a Terraform template in GCP and demonstrated how the Palo Alto Next Generation VM-Series firewall can be deployed via Terraform automation to not only secure traffic throughout your GCP Project, but throughout your Enterprise Google Cloud Infrastructure.

Appendix A

Troubleshooting tips

1. Unable to access the webserver or web page not visible

If the VM-Series firewall is up and accessible but you are unable to access the webserver (or the web page is not visible), then chances are that the startup scripts did not get downloaded from the bootstrap bucket or were corrupted during (or prior to) the upload. Ensure that the files webserver-startup.sh and dbserver-startup.sh are in the bootstrap bucket. If they are extant, replace them with new copies downloaded from the GitHub repository.

2. Bootstrapping not working

If the VM-Series firewall is up and you are able to access the login page, but unable to login using the username/password: paloalto/PaloAlt0@123, then chances are bootstrapping has failed.

There could be several reasons:

a. Corrupt configuration files

Please ensure that the bootstrap.xml and init-cft.txt files mentioned in [Section 4.6](#) are not corrupted.

b. Incorrect bootstrap bucket-name

Another reason for bootstrapping to fail is that the bootstrap bucket name (Parameter: bootstrapbucket) was incorrectly entered in the template file. Please make sure the bucket name created in [Section 4.6](#) is mentioned when launching the template.